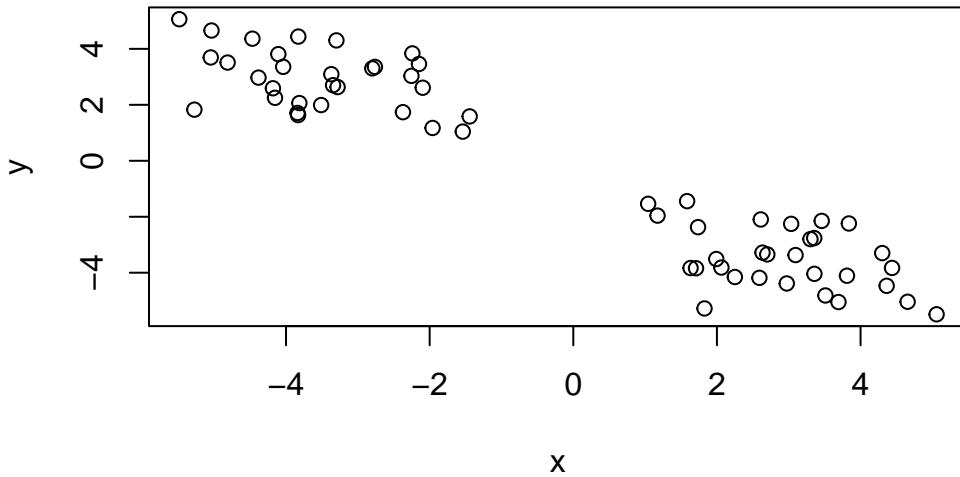


# Class 7: Machine Learning 1

Maddie Maslyar (PID: A69042845)

```
x <- c(rnorm(30, mean = -3), rnorm(30, mean = 3))
y <- rev(x)
z <- cbind(x, y)
```

```
plot(z)
```



The main function in “base” R for K-menas clustering is called ‘kmeans()’

```
k <- kmeans(z, centers = 2, nstart = 20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

### Cluster means:

```

          x           y
1 -3.491002  2.926491
2  2.926491 -3.491002

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 70.90298 70.90298
```

(between\_SS / total\_SS = 89.7 %)

## Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. How big are the clusters?

k\$size

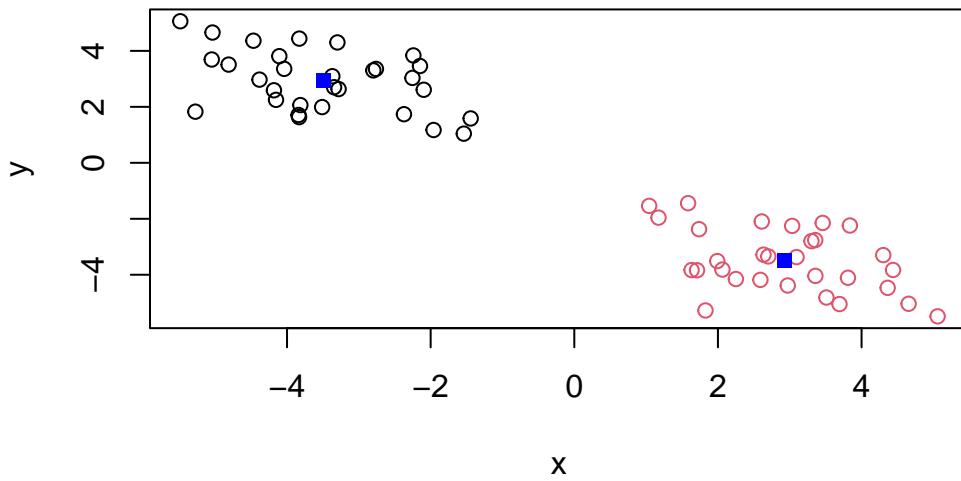
[1] 30 30

What clusters do my data reside within?

k\$cluster

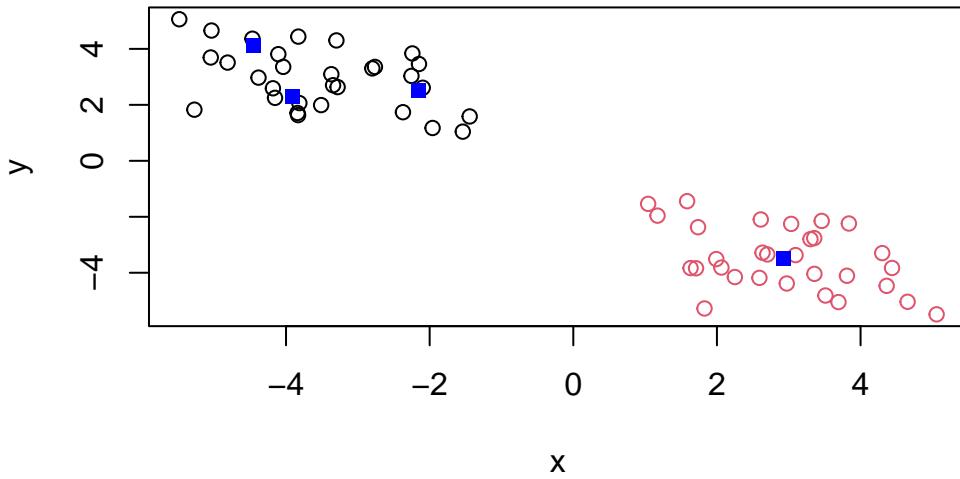
Make a plot of our data colored by cluster assignment:

```
plot(z, col = k$cluster)
points(k$centers, col = "blue", pch = 15)
```



Q. Cluster with k-means into 4 clusters and plot the results.

```
k4 <- kmeans(z, centers = 4)
plot(z, col = k$cluster)
points(k4$centers, col = "blue", pch = 15)
```



Q. Run kmeans with centers (values of k) equal 1 to 6.

```
k$tot.withinss
```

```
[1] 141.806
```

```
k1 <- kmeans(x, centers = 1)$tot.withinss
k2 <- kmeans(x, centers = 2)$tot.withinss
k3 <- kmeans(x, centers = 3)$tot.withinss
k4 <- kmeans(x, centers = 4)$tot.withinss
k5 <- kmeans(x, centers = 5)$tot.withinss
k6 <- kmeans(x, centers = 6)$tot.withinss

ans <- c(k1, k2, k3, k4, k5, k6)
```

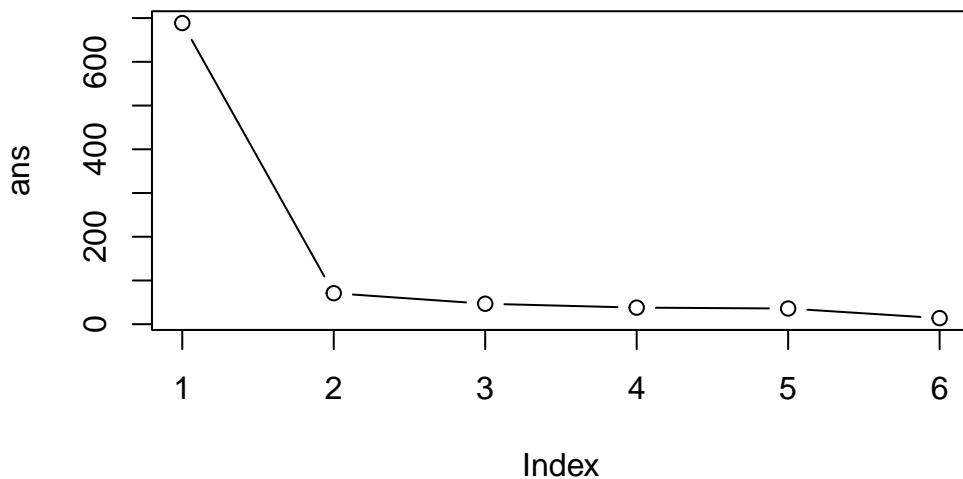
Using a loop:

```
ans <- NULL
for(i in 1:6) {
  ans <- c(ans, kmeans(x, centers = i)$tot.withinss)
}
ans
```

```
[1] 688.66618 70.90298 46.77753 37.91591 35.89244 13.85146
```

Make a “scree-plot”

```
plot(ans, typ = "b")
```



## Hierarchical Clustering

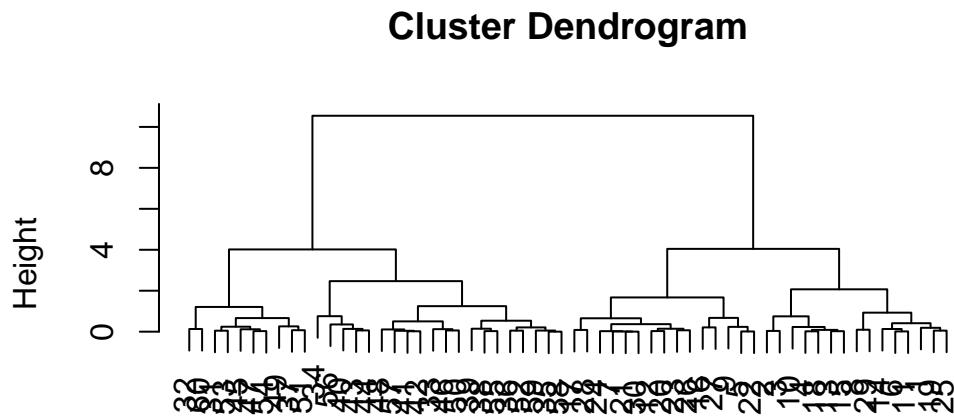
The main function in “base” R for this is called ‘`hclust()`’

```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:  
`hclust(d = d)`

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

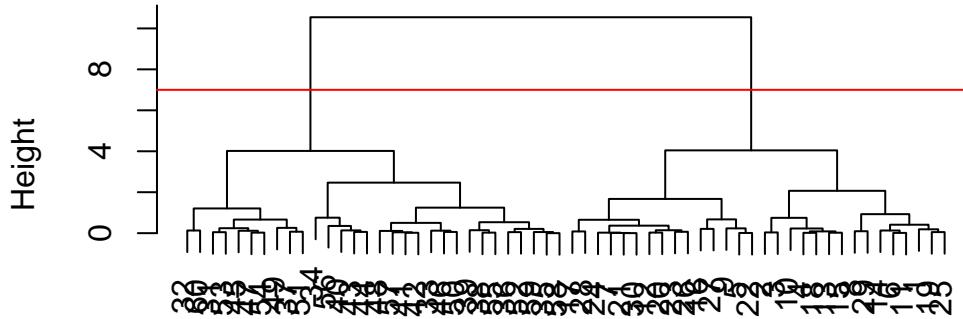
```
plot(hc)
```



d  
hclust (\*, "complete")

```
plot(hc)
abline(h = 7, col = "red")
```

## Cluster Dendrogram



```
d  
hclust (*, "complete")
```

```
library(pheatmap)
```

Class 7 Lab Input information

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
#Set row names to first column of data frame  
  
x <- read.csv(url, row.names = 1)  
head(x)
```

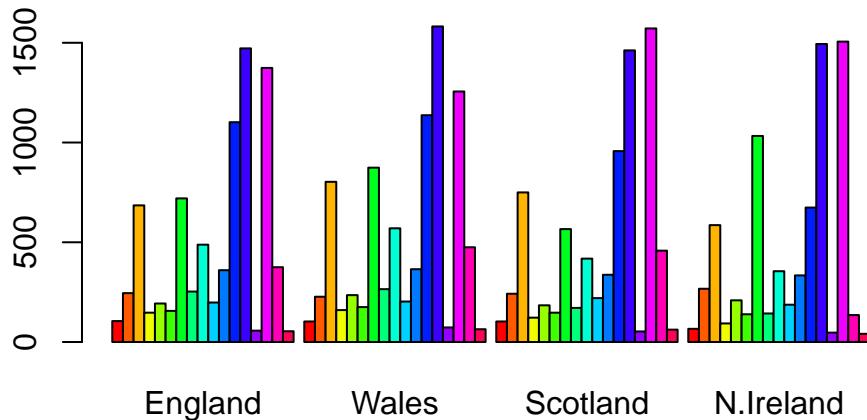
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
dim(x)
```

```
[1] 17 4
```

- A1. There are 17 rows and four columns, as identified by using dim().
- Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?
- A2. Setting the row names equal to 1 prevents the code block from amputating a column from our data table every time the code is run (as was the case for  $x <- x[-1]$ )

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above barplot() function results in the following plot?

A3. Setting beside = FALSE stacks the column values on top of each other.

Q4: Changing what optional argument in the above ggplot() code results in a stacked barplot figure?

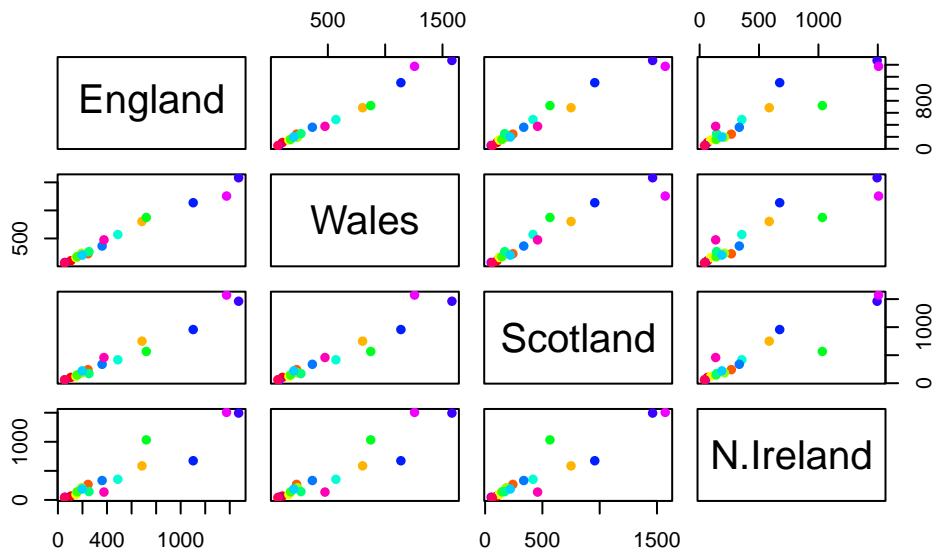
A4. geom\_col creates a stacked bar chart.

### Pairs plots and heatmaps

Scatterplot matrices (“pairs plots”) can be useful for small datasets.

Q5: We can use the pairs() function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

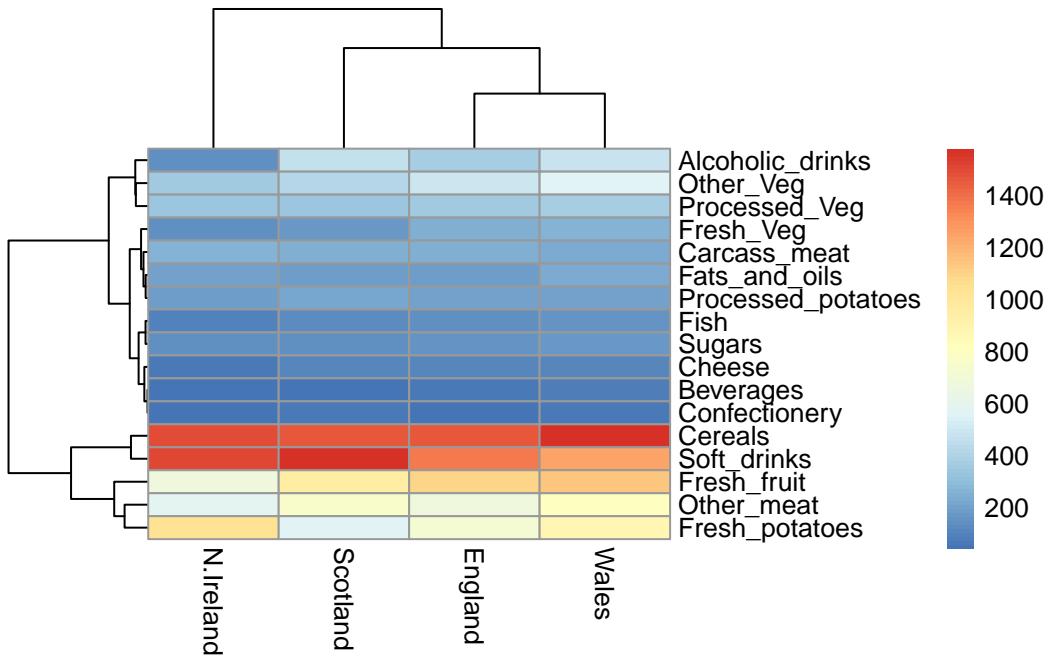
```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



A5. Each point indicates a food that is mapped based on consumption in one country versus another. A point lying along the horizontal line indicates equal consumption in the compared countries.

Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

```
library(pheatmap)
pheatmap( as.matrix(x) )
```



A6. Scotland, England, and Wales cluster, indicating similar food consumption. Northern Ireland has the most distinct clustering. Potatoes and alcoholic drinks appear to be two major differences.

### PCA to the rescue

The main function in “base” R for PCA is called ‘prcomp()’.

As we want to do PCA on the food data for the different countries we will want the foods in the columns.

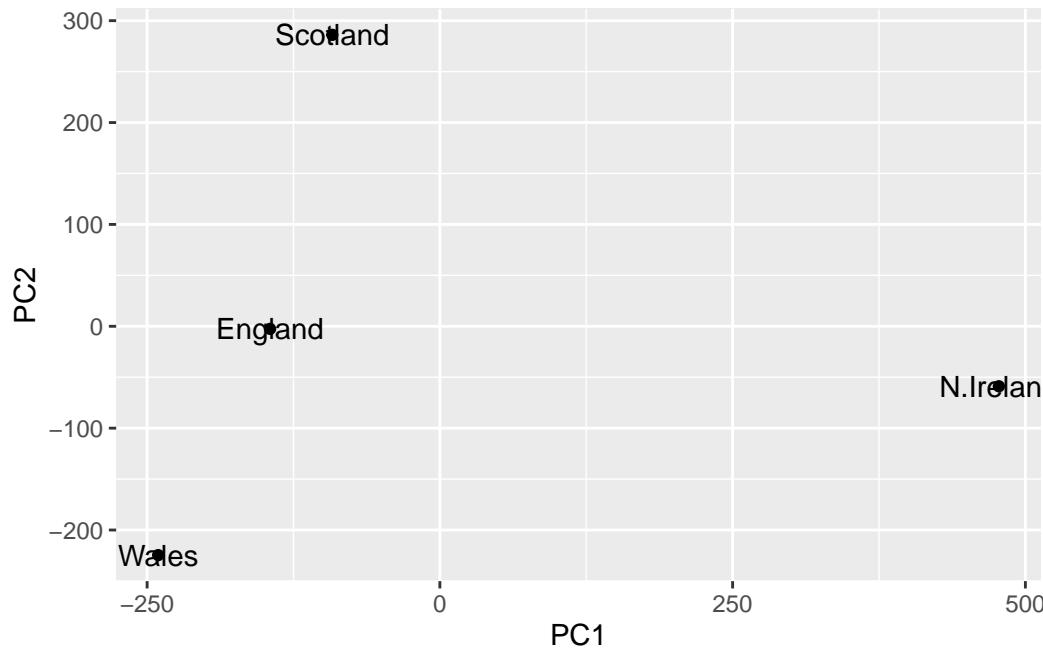
```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

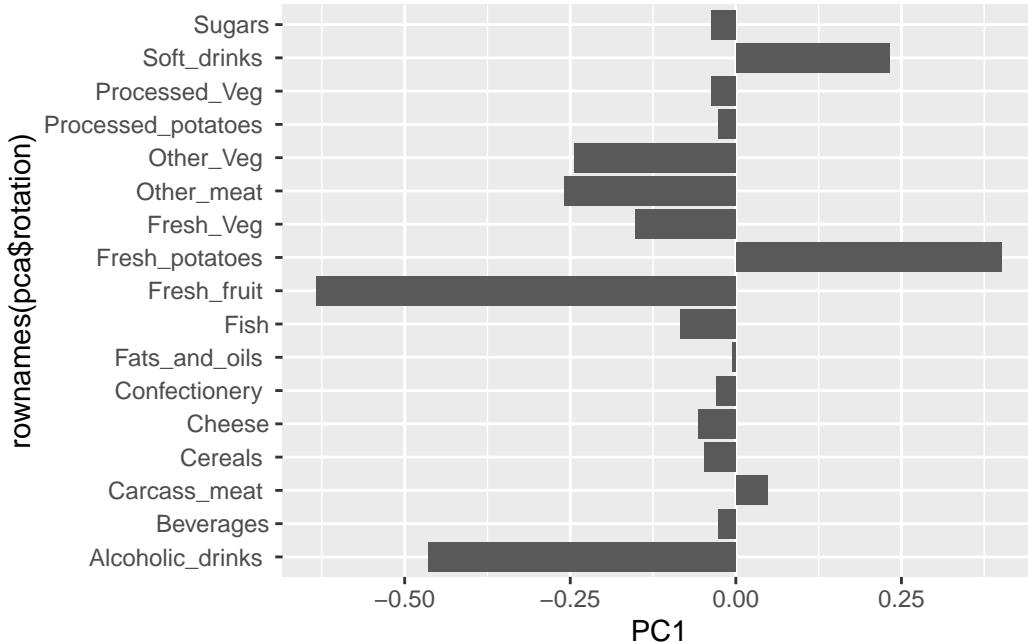
Our result object is called ‘pca’ and it has a ‘\$’ component that we will look at first.

```
library(ggplot2)
ggplot(pca$x) +
  aes(PC1, PC2, label = rownames(pca$x)) +
  geom_point() +
  geom_text()
```



Another major result out of PCA is the so-called “variable loadings” or \$rotation that tells us how the original variables (foods) contribute to PCs (i.e. our new axis).

```
ggplot(pca$rotation) +
  aes(PC1, rownames(pca$rotation)) +
  geom_col()
```

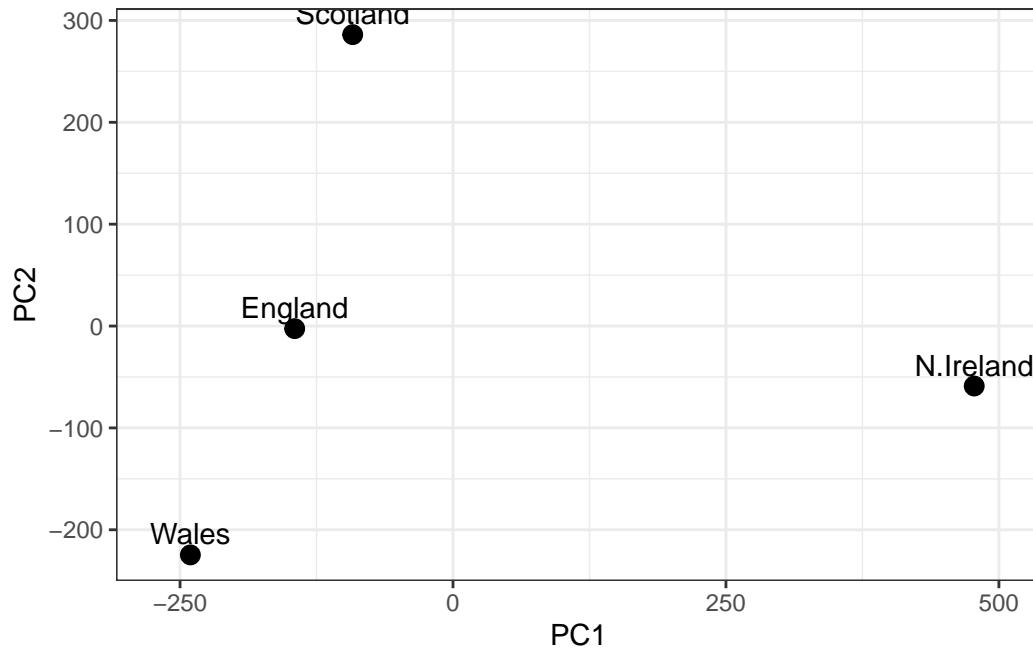


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

A7.

```
# Create a data frame for plotting
df <- as.data.frame(pca$x)
df$Country <- rownames(df)

# Plot PC1 vs PC2 with ggplot
ggplot(pca$x) +
  aes(x = df$PC1, y = df$PC2, label = rownames(pca$x)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

A8.

```
ggplot(pca$x, aes(x = df$PC1, y = df$PC2, label = rownames(pca$x), color = rownames(pca$x)))
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw() +
  scale_color_manual(values = c("yellow", "green", "blue", "red"))
```

